

# Homework: KNN and SVM Implementation

**Course:** Artificial Intelligence

**Department of** Computer Science

**University of** South Dakota

**Supervisor:** *Dr. Lina Chato*

**Semester:** *Spring 2025*

---



## Student Information

Student Name	Course Number & Section
Narayan Bhattarai	CSC 547 UT3
Alex Kagozi	CSC 547 UT3

---

## A. Wine data classification using SVM

Collab and Github Links



[Open in Colab](#)



[View on GitHub](#)

### 1. Import the libraries

```
In [1]: ## Load dependencies
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.datasets import load_wine
from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.svm import SVC
from sklearn.preprocessing import label_binarize, StandardScaler
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
from itertools import cycle
import seaborn as sns
```

### 2. Load the dataset

```
In [2]: # Load dataset
data = load_wine()
```

```
X, y = data.data, data.target
feature_names = data.feature_names
```

3. Show the unique targets and features:

*The dataset has 13 features and 3 targets.*

```
In [3]: ## show the target names and the number of unique features and target
print(data.target_names)
print(len(np.unique(y)))
print(len(feature_names))
print(feature_names)
```

```
['class_0' 'class_1' 'class_2']
```

```
3
```

```
13
```

```
['alcohol', 'malic_acid', 'ash', 'alcalinity_of_ash', 'magnesium', 'total_phenols', 'flavanoids', 'nonflavanoid_phenols', 'proanthocyanins', 'color_intensity', 'hue', 'od280/od315_of_diluted_wines', 'proline']
```

4. Display the features and labels in tabular format:

```
In [4]: # Create a DataFrame with features and target
df = pd.DataFrame(X, columns=feature_names)
df['target'] = y

# Display the first 5 rows
print(df.head(5))
```

	alcohol	malic_acid	ash	alcalinity_of_ash	magnesium	total_phenols	\
0	14.23	1.71	2.43	15.6	127.0	2.80	
1	13.20	1.78	2.14	11.2	100.0	2.65	
2	13.16	2.36	2.67	18.6	101.0	2.80	
3	14.37	1.95	2.50	16.8	113.0	3.85	
4	13.24	2.59	2.87	21.0	118.0	2.80	

	flavanoids	nonflavanoid_phenols	proanthocyanins	color_intensity	hue	\
0	3.06		0.28	2.29	5.64	1.04
1	2.76		0.26	1.28	4.38	1.05
2	3.24		0.30	2.81	5.68	1.03
3	3.49		0.24	2.18	7.80	0.86
4	2.69		0.39	1.82	4.32	1.04

	od280/od315_of_diluted_wines	proline	target
0	3.92	1065.0	0
1	3.40	1050.0	0
2	3.17	1185.0	0
3	3.45	1480.0	0
4	2.93	735.0	0

5. Split the dataset (70 % for training and 30% for testing):

```
In [5]: # Split into train (70%) and test (30%)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, ran
```

6. Standardize the features:

```
In [6]: # Standardize features
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
```

7. Train the data using an SVM linear Kernel:

```
In [7]: # Train Linear SVM Model
svm_model = SVC(kernel='linear', probability=True, random_state=42)
svm_model.fit(X_train, y_train)
```

```
Out[7]: SVC
SVC(kernel='linear', probability=True, random_state=42)
```

8. Run Predictions on Test data and get the accuracy score:

```
In [8]: # Predict on test data
y_pred = svm_model.predict(X_test)

# Compute accuracy
accuracy = accuracy_score(y_test, y_pred)
print(f"Model Accuracy: {accuracy:.2f}")
```


Model Accuracy: 0.96

9. Plot a confusion matrix for the dataset:

```
In [9]: conf_matrix = confusion_matrix(y_test, y_pred)
print("Confusion Matrix:")
print(conf_matrix)
# Plot Confusion Matrix
plt.figure(figsize=(6, 5))
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues', xticklabels=data
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.title('Confusion Matrix')
plt.show()
```

Confusion Matrix:

```
[[18  0  0]
 [ 1 20  0]
 [ 0  1 14]]
```



10. Print a classification report:

```
In [10]: # Classification Report
print("Classification Report:")
print(classification_report(y_test, y_pred))

# Compute ROC curve and AUC
y_test_binarized = label_binarize(y_test, classes=[0, 1, 2])
y_score = svm_model.decision_function(X_test)
```

Classification Report:


	precision	recall	f1-score	support
0	0.95	1.00	0.97	18
1	0.95	0.95	0.95	21
2	1.00	0.93	0.97	15
accuracy			0.96	54
macro avg	0.97	0.96	0.96	54
weighted avg	0.96	0.96	0.96	54

11. Plot a ROC Curve for each class:

```
In [11]: n_classes = len(np.unique(y))

plt.figure(figsize=(8, 6))
colors = cycle(['blue', 'red', 'green'])
for i, color in zip(range(n_classes), colors):
    fpr, tpr, _ = roc_curve(y_test_binarized[:, i], y_score[:, i])
    roc_auc = auc(fpr, tpr)
    plt.plot(fpr, tpr, color=color, lw=2, label=f'Class {i} (AUC = {roc_auc:

plt.plot([0, 1], [0, 1], 'k--', lw=2)
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve')
plt.legend(loc='lower right')
plt.show()
```



## 12. Test model generalization using 5-Fold Cross-Validation

```
In [12]: # Perform 5-Fold Cross-Validation
cv_scores = cross_val_score(svm_model, X_train, y_train, cv=5)
print(f"Cross-Validation Scores: {cv_scores}")
print(f"Mean Cross-Validation Accuracy: {cv_scores.mean():.2f}")
```

Cross-Validation Scores: [0.96      0.96      1.      1.      0.9583  
3333]

Mean Cross-Validation Accuracy: 0.98

## B. Predict Diabetes progression using KNN regression

Collab and Github Links



[View on GitHub](#)

### 1. Import the libraries

```
In [13]: ## Load dependencies
import numpy as np
import matplotlib.pyplot as plt
from sklearn.datasets import load_diabetes
from sklearn.model_selection import train_test_split, cross_val_score
# from sklearn.svm import SVC
from sklearn.neighbors import KNeighborsRegressor
from sklearn.preprocessing import label_binarize, StandardScaler
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
from itertools import cycle
import seaborn as sns
```

### 2. Load the dataset

```
In [19]: # Load dataset
# Load dataset
data = load_diabetes()
X, y = data.data, data.target
feature_names = data.feature_names
```

### 3. Show the unique targets and features:

```
In [20]: ## show the target names and the number of unique features and target
print("Unique target values (rounded):", len(np.unique(np.round(y))))
print("Number of features:", len(feature_names))
print("Feature names:", feature_names)
```

Unique target values (rounded): 214

Number of features: 10

Feature names: ['age', 'sex', 'bmi', 'bp', 's1', 's2', 's3', 's4', 's5', 's6']

### 4. Display the features and labels in tabular format:

```
In [21]: # Create DataFrame with features
df = pd.DataFrame(X, columns=feature_names)
# Add target column
```

```
df['target'] = y
# Display the first 10 rows
print(df.head(10))
```

	age	sex	bmi	bp	s1	s2	s3	\
0	0.038076	0.050680	0.061696	0.021872	-0.044223	-0.034821	-0.043401	
1	-0.001882	-0.044642	-0.051474	-0.026328	-0.008449	-0.019163	0.074412	
2	0.085299	0.050680	0.044451	-0.005670	-0.045599	-0.034194	-0.032356	
3	-0.089063	-0.044642	-0.011595	-0.036656	0.012191	0.024991	-0.036038	
4	0.005383	-0.044642	-0.036385	0.021872	0.003935	0.015596	0.008142	
5	-0.092695	-0.044642	-0.040696	-0.019442	-0.068991	-0.079288	0.041277	
6	-0.045472	0.050680	-0.047163	-0.015999	-0.040096	-0.024800	0.000779	
7	0.063504	0.050680	-0.001895	0.066629	0.090620	0.108914	0.022869	
8	0.041708	0.050680	0.061696	-0.040099	-0.013953	0.006202	-0.028674	
9	-0.070900	-0.044642	0.039062	-0.033213	-0.012577	-0.034508	-0.024993	

	s4	s5	s6	target
0	-0.002592	0.019907	-0.017646	151.0
1	-0.039493	-0.068332	-0.092204	75.0
2	-0.002592	0.002861	-0.025930	141.0
3	0.034309	0.022688	-0.009362	206.0
4	-0.002592	-0.031988	-0.046641	135.0
5	-0.076395	-0.041176	-0.096346	97.0
6	-0.039493	-0.062917	-0.038357	138.0
7	0.017703	-0.035816	0.003064	63.0
8	-0.002592	-0.014960	0.011349	110.0
9	-0.002592	0.067737	-0.013504	310.0

5. Split the dataset (70 % for training and 30% for testing):

```
In [22]: # Split into train (70%) and test (30%)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
print("Training set size:", X_train.shape[0])
print("Test set size:", X_test.shape[0])
```

Training set size: 309  
Test set size: 133

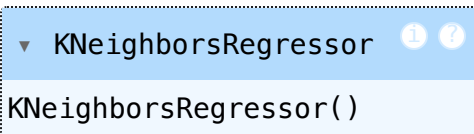
6. Standardize the features

```
In [23]: # Standardize features
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
```

7. Create a KNN Regression model and fit with train data

```
In [24]: # Train KNN Regression Model (K=5)
knn_model = KNeighborsRegressor(n_neighbors=5)
knn_model.fit(X_train, y_train)
```

Out [24]:

A screenshot of a Jupyter Notebook cell output. It shows a dropdown menu with 'KNeighborsRegressor' selected. Below the dropdown, the text 'KNeighborsRegressor()' is displayed. There are two small circular icons to the right of the dropdown menu.

```
▼ KNeighborsRegressor  
KNeighborsRegressor()
```

8. Run prediction on test data

```
In [25]: # Predict on test data  
y_pred = knn_model.predict(X_test)
```

9. Compute the Mean Squared Error


```
In [26]: # Compute Mean Squared Error  
mse = mean_squared_error(y_test, y_pred)  
print(f"Mean Squared Error: {mse:.2f}")
```

Mean Squared Error: 3277.74

10. Plot a Residual Scatter Plot


```
In [27]: # Residual Plot  
residuals = y_test - y_pred  
plt.figure(figsize=(6, 5))  
sns.scatterplot(x=y_pred, y=residuals, alpha=0.7)  
plt.axhline(y=0, color='r', linestyle='--')  
plt.xlabel('Predicted Values')  
plt.ylabel('Residuals')  
plt.title('Residual Plot')  
plt.show()
```





11. Display Predicted vs Actual Values on test set

```
In [28]: # Predicted vs. Actual Plot
plt.figure(figsize=(6, 5))
sns.scatterplot(x=y_test, y=y_pred, alpha=0.7)
plt.plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()], 'r--')
plt.xlabel('Actual Values')
plt.ylabel('Predicted Values')
plt.title('Predicted vs. Actual Values')
plt.show()
```



12. Calculate the Root Mean Squared Error

```
In [29]: # Root Mean Squared Error
rmse = np.sqrt(mse)
print(f"Root Mean Squared Error: {rmse:.2f}")
```

Root Mean Squared Error: 57.25

13. Perform 5-Fold Cross-Validation


```
In [30]: # Perform 5-Fold Cross-Validation
cv_scores = cross_val_score(knn_model, X, y, cv=5, scoring='neg_mean_squared_error')
cv_mse = -np.mean(cv_scores)
print(f"Cross-Validation MSE Scores: {-cv_scores}")
print(f"Mean Cross-Validation MSE: {cv_mse:.2f}")
```

Cross-Validation MSE Scores: [3167.19191011 4022.18067416 3556.56681818 3551.07136364 3807.57 ]  
Mean Cross-Validation MSE: 3620.92

14. Plot a Histogram of Residuals

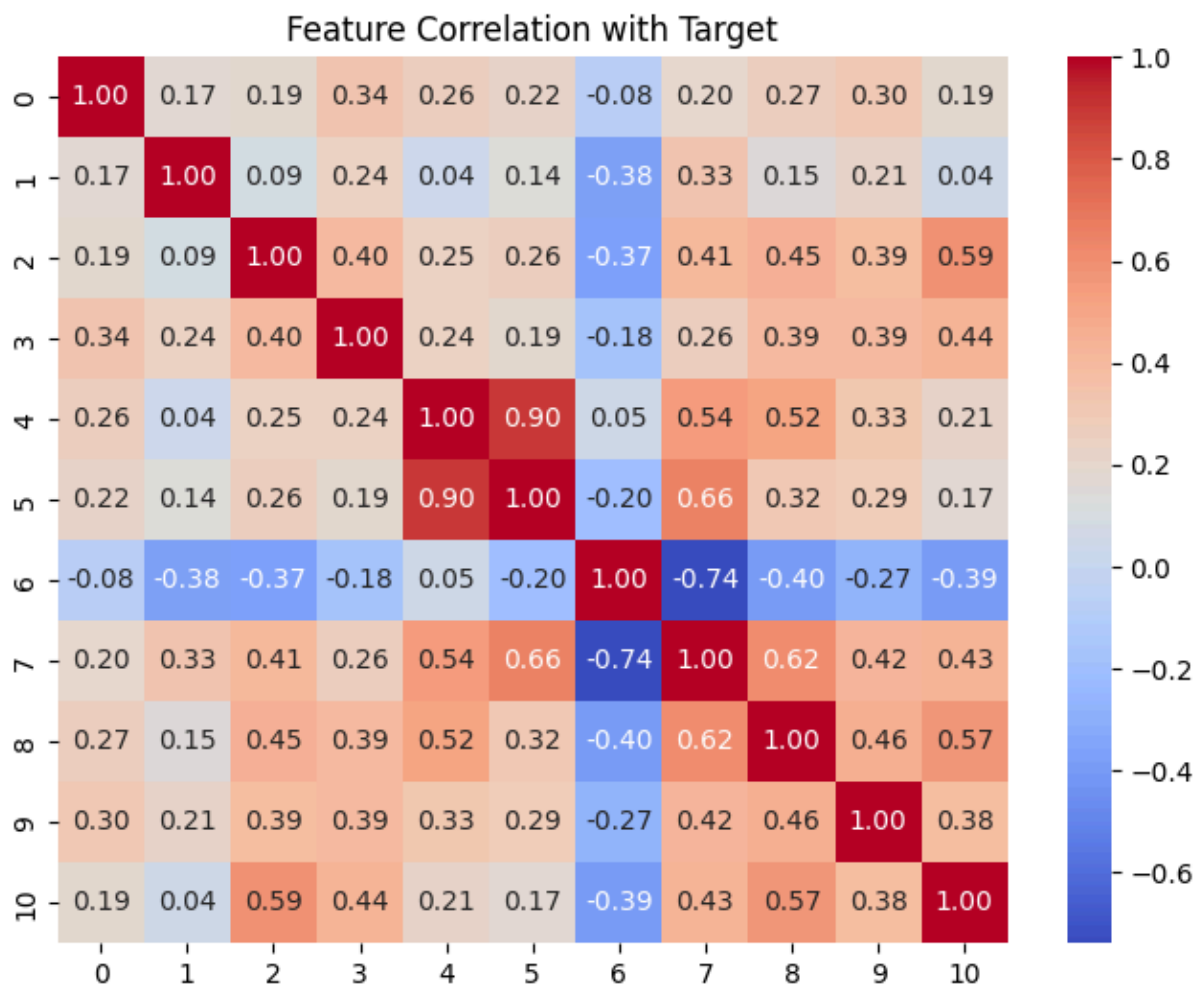
```
In [31]: #Histogram of Residuals
plt.figure(figsize=(6, 5))
sns.histplot(residuals, bins=30, kde=True)
```

```
plt.xlabel('Residuals')  
plt.ylabel('Frequency')  
plt.title('Distribution of Residuals')  
plt.show()
```



15. Show colleration between features and the target

```
In [32]: plt.figure(figsize=(8, 6))  
sns.heatmap(np.corrcoef(X.T, y), annot=True, fmt=".2f", cmap="coolwarm")  
plt.title("Feature Correlation with Target")  
plt.show()
```



In [ ]: